

Introduction to Machine Learning (by Implementation)

Lecture 8: Backpropagation

Ian J. Watson

University of Seoul


University of Seoul Graduate Course 2019



KRF KOREA RESEARCH FELLOWSHIP
해외 우수신진연구자 유치사업

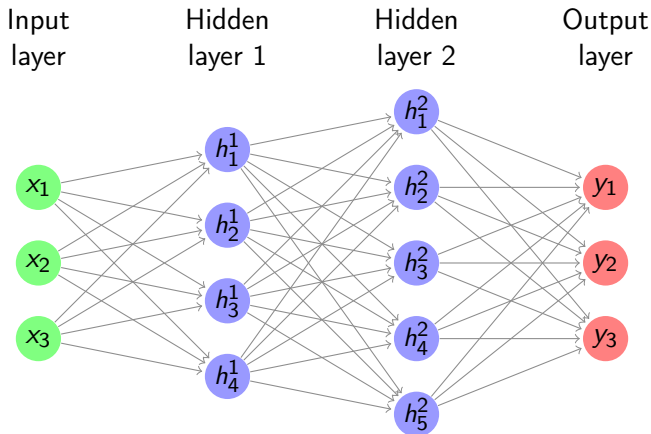


The Feed-Forward Neural Network





- A small feed-forward neural network
 - $y = f(x_1, x_2) = \sigma(-1 + 2x_1 + 2x_2) + \sigma(1 - 2x_1 - 2x_2)$
- Decompose the function into:
 - the *input layer* of \hat{x} ,
 - the *hidden layer* which calculates $h_i = \beta_i \cdot x$ then passes it through the *activation function* σ , (called "sigmoid" in NN terms)
 - as in logistic, there is an extra β_0 , called the *bias*, which controls how big the input into the node must be to activate
 - the *output layer* which sums the results of the hidden layer and gives y
 - $y = \sigma(0 + 1 \cdot h_1 + 1 \cdot h_2)$

Feed-Forward Neural Network



- We can even have several hidden layers
 - The previous layer acts the same as an *input layer* to the next layer
- We call each node in the network a *neuron*
 - At each neuron, the output of the node is $\sigma(\sum \text{weighted node inputs} + \text{bias})$


Training a Neural Network



- What does it mean to train a neural network?
- Consider the XNOR network from last week
- There we set by hand, but could try to "train" the network
- Start with random weights and biases, reduce the loss function
$$C(x, y|w, b) = \sum_i |y_i^{\text{true}} - y(x_i)|^2$$
 where i ranges over our 4 samples (x_i, y_i) and $y(x_i)$ is the network output
 - And, of course, the way we've seen to do this is using *gradient descent*


Gradient Descent on a Neural Network

- Consider running gradient descent on a neural network
- For some particular weight, w_{jk}^l , we want to find $\frac{\partial L}{\partial w_{jk}^l}$
- We could look at this and say, it's big, complicated, let's use our gradient estimator: $\frac{\partial L}{\partial w_{jk}^l} = \frac{L(w_{jk}^l + \Delta) - L(w_{jk}^l)}{\Delta}$ for some small Δ
- But in large networks, we can have millions of nodes: each evaluation of L requires one forward pass through the network, and we need two (at least) for each weight/bias
 - **This means millions of forward passes through the network for a single update**
- And remember, our stochastic algorithm used an update *per known datapoint*
- We need a better way ...



- Of course, the network has a very particular structure: series of evaluations passed from one layer to another, sums inside functions
- Some notation:
 - We have a network of L layers [input layer 0, output layer L]
 - j 'th node on the l 'th layer have output
$$a_j^l = \sigma(z_j^l) = \sigma(\sum_k w_{jk}^l a_k^{l-1} + b_j^l)$$
 - So, the output of the network is a_j^L
 - and the the inputs $x_j = a_j^0$

Backpropagation



- It turns out (from the chain rule), that the gradients can be calculated very simply with one forward pass, and one backward pass propagating the derivatives (hence *backpropagation*)
- Imagine we sit at node a_j^l and we want to find the derivative of w_{jk}^l
 - $\frac{\partial L}{\partial w_{jk}^l} = a_k^{l-1} \delta_j^l$, $\frac{\partial L}{\partial b_j^l} = \delta_j^l$
 - $\delta_j^l = \sigma'(z_j^l) \sum_{k'} w_{k'j}^{l+1} \delta_{k'}^{l+1}$
- That is, the derivative is a product of the activation in a_k^{l-1} and the weighted sum of derivatives coming from the outputs $\delta_{k'}^{l+1}$,
 - Notice that the δ_j^l we calculate on this layer will then be used when setting weights on layer $l + 1$

Backpropagation at the Output Layer

- $\delta_j^l = \sigma'(z_j^l) \sum_{k'} w_{k'j}^{l+1} \delta_{k'}^{l+1}$ can be thought of as the error of the node (look closely on previous page, all w_{jk}^l use the same δ_j^l)
- So, where does it originally come from?
- Well, at the final layer there is no δ^{L+1} to be able to use, so this is our starting point by considering the cost function
- $C = \frac{1}{2} \sum_j (y_j - a_j^L)^2 = \frac{1}{2} \sum_j (y_j - \sigma(z_j^L))^2 = \frac{1}{2} \sum_j (y_j - \sigma(\sum_k w_{jk}^L a_k^{L-1} + b_j^L))^2$
 - Think of the chain rule operating on the expanding piece at each step
- $\frac{\partial C}{\partial w_{jk}^L} = (a_j^L - y_j) \sigma'(z_j^L) a_k^{L-1} = a_k^{L-1} \delta_j^L$, $\frac{\partial C}{\partial b_j^L} = (a_j^L - y_j) \sigma'(z_j^L) = \delta_j^L$
- So, $\delta_j^L = (y_j - a_j^L) \sigma'(z_j^L)$ is our starting point for the backpropagation
 - Use it to set the weights on layer L , then go back a layer, use it as input to find δ_j^{L-1} and then set the weights on layer $L - 1$ and so on
- Notice in the derivation, there was no particular property of σ used other than the fact that we can differentiate it
 - Implies that any activation function will work for backpropagation

Backpropagation Equations and Operation

- $\delta_j^L = (a_j^L - y_j)\sigma'(z_j^L)$
- $\delta_j^l = \sigma'(z_j^l) \sum_{k'} w_{k'j}^{l+1} \delta_{k'}^{l+1}$
- $\frac{\partial L}{\partial w_{jk}^l} = a_k^{l-1} \delta_j^l$
- $\frac{\partial L}{\partial b_j^l} = \delta_j^l$
- TODO: WRITE OUT DEFNS of a vs z
- In the same way that the a_j^l are wrapping up the weighted sums and activations of the layers feeding forward, the δ_j^l wrap up the partial derivatives of the chain rule which must be expanding from the cost function
 - Hopefully, you can see how the proof for the transfer to previous layer would work by running further expansions of a_k^{l-1} on the previous page
- We calculate the a_j^l forward, then calculate the $\frac{\partial C}{\partial w_{jk}^l}, \delta_j^l$ backward
- And then use this to find $\frac{\partial C}{\partial b_j^l}$ and run our SGD

Exercises

- `initialize_weights(n_nodes, initialize_fn=random)`
 - `n_nodes` should be a list of the number of nodes at each layer, including input and output (see the `test_initialize_weights` in `test_neural` for further commentary)
 - Use your `rand.random` function to initialize randomly between 0 and 1
- Should have feedforward from last week, today, lets assume we always use sigmoid activation (so we can use $\sigma'(x) = \sigma(x)(1 - \sigma(x))$)
- `calculate_deltas(network, activations, y)`
 - Calculates the δ_j^l from the previous page
- `batch_update_nn(network, activation, deltas, eta)`
 - Returns the weights after one round of gradient descent updates
 - $w_{jk}^l \rightarrow w_{jk}^l - \eta \frac{\partial C}{\partial w_{jk}^l}$, $b_j^l \rightarrow b_j^l - \eta \frac{\partial C}{\partial b_j^l}$
 - Probably easiest to use `deepcopy` from `copy import deepcopy`, make a copy of the network, then update using indices, rather than trying to make the network as you go

Exercises

- `sgd_nn(x, y, theta0, eta=0.1)`
 - Similar structure as our previous stochastic gradient descent, but uses the functions above to do the updates of the weights on each sample
 - Instead of input functions, assume a sum of squares cost function and use the batch update sequence you've just written `feedforward_`, `calculate_deltas`, `batch_update_nn`
 - It can be useful to save the values of the cost function to monitor how much the network is changing, particularly to try out different `eta`
 - You might find it easier to drop the `n_iterations` and run `n_epochs` (times over dataset) with your own training schedule (`eta` choice)
- Try training a network on our xor problem from last week.
- Hint: use gaussian initialized weights, play with the `alpha` and `n_iterations` hyperparameters. You might need to try it a few times with different starting points to get good convergence
- Try training a network for the Fisher classification problem from two weeks ago
 - Play around with the network architecture (number of layers/nodes)
- Use the `multi_accuracy` and print out your best network and accuracy into `results.txt`